



ENTRY PLAN

90 days to understand, test, and strengthen the proving system.

This plan is intentionally diagnostic before prescriptive. KPMG's actual portfolio, platform architecture, governance, and decision rights are not public in enough detail to pretend otherwise. The first obligation is to learn where the experiment-to-scale system works, where it stalls, and which mechanisms are worth standardizing.

DAYS 0–30

Map the system as it really operates.

Portfolio trace

- Select a representative sample of active, scaled, stalled, and retired AI efforts.
- Trace origin, sponsor, workflow, platform path, evidence, decision owner, adoption state, and current economics.
- Identify recurring failure points rather than isolated project anecdotes.

Stakeholder topology

- Business and domain owners
- AI and platform engineering
- Data architecture and access owners
- Security, privacy, legal, risk, and responsible AI
- Finance / FinOps and investment owners
- Change, learning, enablement, and workflow users
- Cloud, model, and specialist partners

Deliverable: Experiment System Map

A one-page view of demand inflow, framing, POC build, evaluation, risk review, scale decision, operational handoff, adoption ownership, and learning capture—plus a short list of unresolved decision-right ambiguities. No reorganization proposal in month one.

DAYS 31–60

Test a proving discipline on real work.

Choose representative opportunities

Run two or three opportunities with different profiles through the experiment decision card: one high-reuse platform candidate, one domain-specific workflow, and one high-assurance use case.

Define evidence before build

For each, write the work boundary, human authority, data and test requirements, economic ceiling, adoption burden, and explicit scale-or-stop decision.

Stress the proof

Review task quality, groundedness, edge cases, failure recovery, privacy/security posture, latency, run cost, workflow integration, override patterns, and actual operator behavior.

Observe the mechanism

Track which review steps add signal, which duplicate effort, where ownership is unclear, and what evidence different stakeholders need to commit.

Deliverable: Reusable Experiment Contract v0.1

A lightweight, testable brief—not a bureaucracy. It should make POCs faster by resolving uncertainty earlier, clarifying who decides, and making evaluation comparable enough to reuse lessons across work.

DAYS 61–90

Turn learning into institutional memory.

Scale / partner / stop cadence

Create a concise evidence review that ends in an explicit decision, owner, rationale, and next checkpoint rather than indefinite pilot status.

Reusable evaluation pack

Standardize the parts of task evaluation, failure-mode review, cost visibility, and adoption observation that can transfer across experiments.

Capability harvest

Identify common services, patterns, controls, prompts, datasets, connectors, and tests worth turning into shared capabilities.

Operating handoff

Make ownership after scale explicit: support, monitoring, cost, model/provider change, workflow training, review obligations, and retirement.

Deliverable: Proving Ground v1

An evidence-backed operating mechanism fitted to KPMG's actual environment: experiment intake, decision card, bounded proof, stress protocol, scale-or-stop review, and reusable learning capture. The form should follow what the first 60 days reveal.

How I would know the entry plan is working.

CLARITY

Decision rights

Fewer ambiguous POC states and clearer owners for scale, stop, risk acceptance, and operating handoff.

SPEED

Evidence cycle

Shorter time from idea to decision-quality evidence, including faster termination of weak work.

REUSE

Learning carryover

More experiments begin with prior tests, patterns, cost knowledge, controls, and architecture decisions already in hand.