

90 days to understand, test, and strengthen the proving system.

This plan is intentionally diagnostic before prescriptive. KPMG's actual portfolio, platform architecture, governance, and decision rights are not public in enough detail to pretend otherwise. The first obligation is to learn where the experiment-to-scale system works, where it stalls, and which mechanisms are worth standardizing.

OPERATING PREMISE

Innovation compounds when every experiment leaves behind one of three assets: a reusable capability, reusable evidence, or a fast no.

FRAME

Work + outcome

BOUND

Authority + risk

BUILD

Minimum proof

STRESS

Quality + cost + use

COMPOUND

Scale + reuse + learn

01 Days 0-30: Map the system as it really operates.

Build a factual view of how demand becomes an experiment, how evidence is judged, and where accountability changes hands.

Portfolio trace

- Select active, scaled, stalled, and retired AI efforts.
- Trace origin, sponsor, workflow, platform path, evidence, decision owner, adoption state, and economics.
- Identify recurring failure points.

Stakeholder topology

- Business and domain owners
- AI, platform, and data engineering
- Security, privacy, legal, risk, and responsible AI
- Finance / FinOps, change, learning, and partners

Decision-right audit

-

03 Turn learning into institutional memory.

Standardize only what the first 60 days prove is reusable. Make ownership after scale explicit.

Scale / partner / stop cadence
End each review with an explicit decision, owner, rationale, and next checkpoint rather than indefinite pilot status.

Reusable evaluation pack
Standardize transferable task evaluation, failure-mode review, cost visibility, and adoption observation.

Capability harvest
Identify shared services, patterns, controls, prompts, datasets, connectors, and tests worth institutionalizing.

Operating handoff
Make support, monitoring, cost, provider change, workflow training, review obligations, and retirement ownership explicit.

Deliverable: Proving Ground v1
An evidence-backed operating mechanism fitted to KPMG's actual environment: experiment intake, decision card, bounded proof, stress protocol, scale-or-stop review, and reusable learning capture.

Early measures

Clarity | Decision rights
Fewer ambiguous POC states and clearer owners for scale, stop, risk acceptance, and operating handoff.

Speed | Evidence cycle
Shorter time from idea to decision-quality evidence, including faster termination of weak work.

Reuse | Learning carryover
More experiments begin with prior tests, patterns, cost knowledge, controls, and architecture decisions already in hand.

MANAGEMENT TEST
Did the mechanism shorten time to decision-quality evidence while preserving the controls, economics, and professional judgment KPMG requires?

Canonical role source: <https://www.kpmguscareers.com/jobdetail/?jobId=134718-B&srcCat=Job+Boards&specSrc=LinkedIn>

Candidate-created hypothesis: This plan is not represented as KPMG policy. It is a discovery model to test against KPMG's actual portfolio, architecture, governance, and operating constraints.