

The case I would bring into the room.

Prepared from the job posting, current public signals, and verified career evidence. Company-specific recommendations are hypotheses for discovery unless directly sourced.

Actual mandate

Convert GenAI experimentation into compounding enterprise capability. Build the proving discipline between enterprise demand and scaled AI: better problem framing, explicit authority and risk boundaries, enterprise-representative POCs, realistic stress testing, clear scale, partner, or stop decisions, and deliberate reuse of capability and evidence.

Company moment

- KPMG has publicly expanded major AI relationships and platforms across cloud ecosystems.
- Leadership has described a move from experimentation and POCs toward implementation and scaling.
- Recent KPMG-reported research highlights weak visibility into AI operating costs and the importance of AI economics.
- The role itself spans GenAI development, POCs, cross-functional coordination, partners, thought leadership, roadmaps, and business impact.

Core tensions

- Experiment velocity <-> enterprise assurance
- POC novelty <-> reuse and economics
- Platform optionality <-> common architecture and controls
- Technical depth <-> stakeholder translation
- Central innovation <-> domain professional judgment

A proving discipline, not a demo factory.

Candidate thesis

Innovation compounds when every experiment leaves behind one of three assets: a reusable capability, reusable evidence, or a fast no. The GenAI Proving Ground is the operating concept: FRAME the work, BOUND authority and risk, BUILD the minimum enterprise proof, STRESS quality, economics, and adoption, then COMPOUND through scale, reuse, partnership, or disciplined retirement.

Decision architecture

- What changes in the workflow?
- Where does human authority remain?
- Build, buy, partner, or reuse?
- What evidence warrants scale?
- What gets standardized?
- What gets stopped?


Balanced scorecard

- Business movement
- Adoption behavior
- AI task performance and failure modes
- Trust, controls, and auditability
- Run cost and reuse economics
- Time from idea to decision-quality evidence

Proof map

- Vape-Jet: direct AI-first operating transformation across factory, customer, support, engineering issue flow, ERP, collaboration systems, telemetry, and AI-assisted workflows.
- DudeWorth: direct AI-readiness and opportunity frameworks plus agentic workflow patterns.
- Amazon: scaled 24/7 launch and operating mechanisms.
- Compunetics: high-reliability technical program, engineering, manufacturing, quality, and control systems.
- ZeusVu: regulated autonomy, computer-vision concepts, chain of custody, and a 100% safety record with no FAA incidents.

Own the gap. Make the transferability case.



Public sources and discovery boundaries.

Canonical job posting

<https://www.kpmguscareers.com/jobdetail/?jobId=134718-B&srcCat=Job+Boards&specSrc=LinkedIn>

Reuters - KPMG / Google Cloud AI investment and scaling

<https://www.reuters.com/technology/artificial-intelligence/kpmg-spend-100-million-ai-partnership-with-google-cloud-2024-11-20/>

KPMG Global - official visual and editorial system

<https://kpmg.com/xx/en.html>

KPMG Global Careers - people and technology cues

<https://kpmg.com/xx/en/careers.html>

Candidate vision

<https://russelldudek.github.io/KPMG/>

Discovery boundary

The Proving Ground, decision card, scorecard, and 90-day plan are candidate-created hypotheses. They are not represented as KPMG policy. The first interview objective is to test where these ideas fit, where KPMG already has stronger mechanisms, and which constraints are invisible from public sources.

Brand and asset posture

The campaign uses an original blue-led, high-contrast, modular editorial system informed by KPMG official web and careers properties. It does not embed the official KPMG logo or proprietary imagery and does not imply endorsement. Brand provenance is recorded in brand-intelligence.md.