INTERVIEW THESIS BRIEF

The case I would bring into the room.

Prepared from the canonical job posting, current public signals, and verified career evidence. Company-specific recommendations are hypotheses for discovery unless directly sourced.

Actual mandate

Convert GenAI experimentation into compounding enterprise capability. Improve how opportunities are framed, bounded, proven, stressed, scaled, reused, partnered, or stopped.

Company moment

- KPMG is expanding AI investment and cloud/model ecosystem relationships.
- Leadership has described a move from experimentation toward implementation and scale.
- AI operating-cost visibility and value realization are part of the mandate.

Core tensions

- Experiment velocity ↔ enterprise assurance
- POC novelty ↔ reuse and unit economics
- Platform optionality ↔ common controls
- Technical depth ↔ cross-functional translation
- Central innovation ↔ domain judgment

Candidate thesis

Innovation compounds when every experiment leaves behind one of three assets: a reusable capability, reusable evidence, or a fast no.

Operating model: FRAME → BOUND → BUILD → STRESS → COMPOUND.

Decision architecture

- What work changes?
- Where does human authority remain?
- Build, buy, partner, or reuse?
- What evidence warrants scale?
- What gets standardized or stopped?

Balanced scorecard

- Business movement
- Adoption behavior
- Task quality and failure modes
- Trust, controls, and auditability
- Run cost and reuse economics
- Time to decision-quality evidence

Proof map

Vape-Jet: AI-first operating transformation in robotics and machine vision. **DudeWorth:** AI-readiness and agentic workflow design. **Amazon:** scaled 24/7 operating mechanisms. **Compunetics:** high-reliability engineering, manufacturing, quality, and controls. **ZeusVu:** regulated autonomy, computer-vision concepts, and a 100% safety record with no FAA incidents.

Strongest objection

The role asks for deep recent AI/ML experience and hands-on enterprise AI platform development. The verified record does not support claiming a foundation-model research career or unverified Azure AI Foundry / Vertex AI production ownership.

Response strategy

Own the gap. Make the adjacent-evidence case around enterprise POCs, cross-functional engineering, roadmaps, partners, resource priorities, adoption, governance, and business impact. Test platform-specific depth directly in interview.

First 90 days

0–30: map experiments, stakeholders, decision rights, and failure points. **31–60:** test the experiment contract on representative opportunities. **61–90:** publish reusable evaluation patterns, clarify scale/partner/stop cadence, harvest capabilities, and make operating handoff explicit.

Executive questions

1. Where do promising experiments currently stall?
2. Who has final scale/stop authority, and what evidence do they trust?
3. What must be common across platform ecosystems, and what should remain optional?
4. Which workflows justify custom AI capability?
5. How is meaningful adoption distinguished from usage?
6. Who owns AI unit economics after scale?

Sources

- KPMG US Careers — Job 134718
- Reuters — KPMG / Google Cloud AI investment
- KPMG Global — official visual source
- KPMG Global Careers — official careers source